

Programming paradigms for GPU devices



28th Summer School on
Parallel Computing

1-12 July 2019

Sergio Orlandini

s.orlandini@cineca.it

- Matrix Transpose
 - evaluating performance
 - using shared memory for coalesced accesses
 - avoiding bank-conflicts



Matrix Transpose

- Let's implement matrix transpose with the following simple design:
 - out-of-place buffers
 - square matrices with size modulo 32 elements
- This is a memory-bounded kernel
 - no computation on elements
 - just load and stores
- We will use effective bandwidth (GB/s) as a metric to measure the performance of such kernels

Simple Copy Kernel

```
__global__ void copy (float *idata, float *odata, int width, int height)
{
    int xIndex = blockIdx.x * TILE_DIM + threadIdx.x;
    int yIndex = blockIdx.y * TILE_DIM + threadIdx.y;

    int index_in  = width * yIndex + xIndex;
    int index_out = index_in;

    odata[index_out] = idata[index_in];
}
```

- We can use a simple copy kernel as a reference
 - **TILE_DIM** is the size of the square sub-matrix block
 - we map CUDA blocks to sub-matrix blocksmatrix sub-block size == CUDA thread block size

Naive Transpose

```
__global__ void transposeNaive(float *idata, float *odata, int width, int height)
{
    int xIndex = blockIdx.x * TILE_DIM + threadIdx.x;
    int yIndex = blockIdx.y * TILE_DIM + threadIdx.y;

    int index_in  = width * yIndex + xIndex;
    int index_out = height * xIndex + yIndex;

    odata[index_out] = idata[index_in];
}
```

```
attributes(global) subroutine transposeNaive (idata, odata, width, height)
    integer, intent(in), value :: width, height
    real, intent(in) :: idata(width,height)
    real, intent(out) :: odata(height,width)

    i = ( blockIdx%x - 1 ) * TILE_DIM + threadIdx%x
    j = ( blockIdx%y - 1 ) * TILE_DIM + threadIdx%y

    odata(j,i) = idata(i,j)

end subroutine
```

Measuring Performance

```
// take measurements for loop over kernel launches
cudaEventRecord(start);
for (int i=0; i < NUM_REPS; i++)
    kernel<<<grid, threads>>>(
        d_idata, d_odata, width, height);
cudaEventRecord(stop);
cudaEventSynchronize(stop);
float outerTime;
cudaEventElapsedTime(&outerTime, start, stop);

// take measurements for loop inside kernel
cudaEventRecord(start);
kernel<<<grid, threads>>>(
    d_idata, d_odata, width, height, NUM_REPS);
cudaEventRecord(stop);
cudaEventSynchronize(stop);
float innerTime;
cudaEventElapsedTime(&innerTime, start, stop);
```

Effective Bandwidth (GB/s) on 2048x2048 S2050

<i>kernel</i>	<i>Performance [GB/s]</i>
Copy	60.9
Naive	22.4

... what is happening?

Not coalesced Accesses

- All loads from input matrix are coalesced:
 - each *warp* reads a line of contiguous elements
 - 32 float belonging to the same cache line
- yet all stores into transposed matrix are not coalesced:
 - the copy kernel store by lines
 - the naive transpose kernel store by columns
 - threads in a warp write an elements into different segments
 - the matrix stride rules how distant those segments are
- the naive transpose kernel performs 32 different stores per row

Coalesced Transpose



- To avoid non-coalesced store we should store by row:
 - let's fill a tile in shared memory with data to be transposed
 - we don't get any penalty writing elements by columns into shared-memory
 - the transpose operation is now performed in shared-memory
 - once the tile is filled, we write back into global memory by rows

Coalesced Transpose

```
__global__ void transposeCoalesced(float *idata, float *odata, int
width, int height)
{
    __shared__ float tile[TILE_DIM][TILE_DIM];

    int xIndex = blockIdx.x * TILE_DIM + threadIdx.x;
    int yIndex = blockIdx.y * TILE_DIM + threadIdx.y;
    int index_in  = width * yIndex + xIndex;

    xIndex = blockIdx.y * TILE_DIM + threadIdx.x;
    yIndex = blockIdx.x * TILE_DIM + threadIdx.y;
    int index_out = height * yIndex + xIndex;

    tile[threadIdx.y][threadIdx.x] = idata[index_in];

    __syncthreads();

    odata[index_out] = tile[threadIdx.x][threadIdx.y];
}
```

Coalesced Matrix Transpose


Effective Bandwidth (GB/s) on 2048x2048 S2050	
<i>kernel</i>	<i>Performance [GB/s]</i>
Copy	60.9
Naive	22.4
Coalesced	24.1

... mmm, we are still missing something

Avoiding Bank Conflict

- our new coalesced transpose use a shared memory tile of 32x32 float
 - each element resides on successive bank (4-byte)
 - accessing elements with a 32 size stride will fetch them **from the same bank**
 - any read/write access to this tile by column will get a 32 bank conflict
- use the trick! a new tile of 32x33 elements
 - element of the same tile column will reside on different banks
 - no more bank conflicts at all

```
__shared__ float tile[TILE_DIM][TILE_DIM+1];
```



Avoiding Bank Conflict

Effective Bandwidth (GB/s) on 2048x2048 S2050

<i>kernel</i>	<i>Performance [GB/s]</i>
Copy	60.9
Naive	22.4
Coalesced	24.1
no Bank Conflicts	46.6

Rights & Credits

These slides are CINECA 2014 and are released under the Attribution-NonCommercial-NoDerivs (CC BY-NC-ND) Creative Commons license, version 3.0.

Uses not allowed by the above license need explicit, written permission from the copyright owner. For more information see:

<http://creativecommons.org/licenses/by-nc-nd/3.0/>

Slides and examples were authored by:

Isabella Baccarelli, Luca Ferraro, Sergio Orlandini